

COMP90007 Internet Technologies

Network Analysis Assignment

Measuring bandwidth, delay and jitter

Due Date: 5:00pm - May 11, 2018

1 Introduction

The aim of this project is for you to gain practical experience in making measurements in networks and working with common networking command line utilities. You will specifically be involved with the measurement of bandwidth, delay and jitter. These tasks will be similar to those you have performed in the laboratory sessions held in the tutorials. It is recommended that you perform these tasks in a consistent networking environment to reduce the variance in your results.

Important Note: When you run these commands, take a screenshot of the results obtained and place it in the appendix of the document. Reports failing to do so will be penalized. All plots however, should be placed in the main body of your report where you will explain the observations being derived.

2 Measuring the hop count

In this section, we will be observing the number of intermediate hosts in the route taken to communicate with a remote server and its relation to the physical geographical distance.

To count the number of hops taken to reach a destination host, the command `tracert` will be used (or its corresponding equivalent, depending on your operating system). This utility should be pre-installed on your operating system.

The utility can be invoked by launching a command line terminal and typing in the command. An example output of the `tracert` command (on OS X) and `tracert` command (on Windows) is as follows:

```
$ traceroute -nw1 cis.unimelb.edu.au
traceroute to cis.unimelb.edu.au (128.250.37.164), 64 hops max, 52
byte packets
 1  10.0.0.254  533.676 ms  1.063 ms  0.940 ms
 2  58.96.2.205  27.872 ms  28.137 ms  28.293 ms
 3  58.96.2.129  28.647 ms  28.577 ms  28.085 ms
 4  218.100.78.33  28.299 ms  28.469 ms  28.332 ms
 5  202.158.200.9  29.626 ms  28.871 ms  29.841 ms
 6  202.158.210.26  31.320 ms  28.722 ms  29.135 ms
 7  202.158.200.250  29.668 ms  29.096 ms  28.660 ms
 8  * * *
 9  * * *
10  * * *
11  128.250.37.130  957.521 ms  33.475 ms  29.891 ms
12  128.250.37.164  29.940 ms  29.260 ms  30.020 ms
```

In this section of the project, you are interested in the number of hops it takes to reach the destination server. In the example above, the number of hops to reach `cis.unimelb.edu.au` is 12.

Based on the number of measurements you will be taking, there are some useful command line parameters you may wish to take advantage of, to speed up the time it takes to gather results. The help documentation for the `tracert` utility can be accessed by running `man tracert` or `tracert /?` on Windows. You may also wish to investigate shell scripting to automate the collection of results, but this is not required for the project. Any scripts you do choose to write, however, **must be included in the Appendix**.

2.1 Task description (1.5 marks)

Please include all raw measurements in the Appendix.

2.1 What do the command line parameters `-nw1` (equivalently `-d -w 1` on Windows) mean in the example given above and why were they used?

2.2 Determine the hop count for the following hosts given in Table 1.

Table 1: List of public `iperf` hosts

Host	Location
<code>iperf.he.net</code>	USA
<code>bouygues.testdebit.info</code>	France
<code>iperf.comneonext.de</code>	Germany
<code>ikoula.testdebit.info</code>	France
<code>st2.nn.ertelecom.ru</code>	Russia
<code>iperf.biznetnetworks.com</code>	Indonesia
<code>iperf.scottlinux.com</code>	USA
<code>speedtest.serverius.net</code>	Netherlands

Determine the approximate geographical distance for the above hosts and plot the hop count versus the approximate geographical distance from Melbourne. Do you observe a correlation? Why? Why not?

You may use any scientific computing package or spreadsheet software to do your plotting, **for example, Microsoft Excel**.

Note: These servers are public servers and are not maintained by the University of Melbourne, hence they are likely to go down at any point in time. Based on past experience, it would be advisable to conduct your tests on these as soon as possible (ASAP) rather than leave it till the end of the time frame because there is a high probability that these servers might not be available then (no longer work). If these servers stop responding then please visit the link: <https://iperf.fr/iperf-servers.php> and find your own servers (anything that is responding) or feel free to find any public `iperf` server from the internet, including deploying your own `iperf` server.

3 Measuring delay and jitter

In this section, you will be measuring the delay and jitter of different hosts located in different geographical locations.

We will be using the `ping` utility, to measure the round-trip delay of packets. The `ping` utility should be pre-installed on all major operating systems. The standard deviation of the round-trip delay time will be taken as the value for *jitter* for this project.

The standard deviation measures the variation in a set of data. It is defined as the square root of the variance and **is expressed as follows** :

$$\sigma = \sqrt{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2}$$

where \bar{x} is the mean of the set of data.

A sample output of the ping utility is shown below, but this output will vary depending on your operating system.

```
$ ping unimelb.edu.au
PING unimelb.edu.au (172.22.44.10): 56 data bytes
64 bytes from 172.22.44.10: icmp_seq=0 ttl=124 time=3.364 ms
64 bytes from 172.22.44.10: icmp_seq=1 ttl=124 time=3.416 ms
64 bytes from 172.22.44.10: icmp_seq=2 ttl=124 time=3.730 ms
^C
--- unimelb.edu.au ping statistics ---
3 packets transmitted, 3 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 3.364/3.503/3.730/0.162 ms
```

In the output above, various statistics, including the mean and standard deviation, were calculated for you by the utility. It is expected that you record all the values of delay and perform your own calculations to find the mean values and their standard deviation, to confirm the result.

3.1 Task description (3 marks)

Please include all raw measurements in the Appendix.

3.1 What are the user-facing implications of high delay and high jitter in networking applications? Give specific examples on applications that are sensitive to both high delay and high jitter. Why is this the case?

3.2 Measure the round-trip delay for the following hosts. Make **three** delay measurements (run this command 3 times) of each host and find the average round-trip delay and jitter by calculating the standard deviation, for each host in Table 1.

For each of the above hosts, plot the average round-trip delay versus the approximate physical geographical distance to the server. Do the same with the jitter.

For finding out the physical geographical distance you may use any tool or application available online, for example, you may use a combination of: <https://db-ip.com> and <https://www.freemaptools.com/how-far-is-it-between.htm> or <https://www.site24x7.com/find-website-location.html> and <https://www.distancecalculator.net/> or anything of your choice. However, do make sure to document it and provide the appropriate reference to that application/ tool/ software used.

3.3 From the two plots above, do you observe any correlation between delay and jitter as a function of distance? Why? Why not? Explain your results with reference to the network environment in which you were collecting your results.

4 Measuring the bandwidth-delay product

In this section, we will be measuring the bandwidth of different hosts in order to determine the bandwidth-delay product, using the results from the previous sections.

The utility that will be used to perform bandwidth measurements will be the `iperf` utility. This command line utility is available for download for all operating systems from <https://iperf.fr>. Alternatively you may choose to use the package manager for your relevant operating system.

There are two modes of operation in `iperf`. The server mode will host a server which will listen to incoming requests from a client. An `iperf` instance running in client mode will connect to the server, and packets will be exchanged and timed between the two hosts to calculate the bandwidth. In this project, we will be running `iperf` in client mode.

A sample output of `iperf` in client mode is shown below, noting the `-c` flag to designate operating in client mode.

```
$ iperf -c iperf.eenet.ee
```

```
-----  
Client connecting to iperf.eenet.ee, TCP port 5001  
TCP window size: 129 KByte (default)  
-----
```

```
[ 5] local 10.0.0.5 port 51878 connected with 193.40.100.7 port 5001  
[ ID] Interval      Transfer      Bandwidth  
[ 5]  0.0-106.7 sec  128 KBytes   9.83 Kbits/sec
```

(Note: speedtest.serverius.net server) For speedtest.serverius.net, we need to use the port 5002, so the command will be: "iperf3 -c speedtest.serverius.net -p 5002" for this host. Also, **some iperf servers respond to iperf2, rest to iperf3** so you might want to use both iperf2 and iperf3 to verify if the server is responsive.

4.1 Task description (5.5 marks)

Please include all raw measurements in the Appendix.

- 4.1 What does the bandwidth-delay product tell us about the networks. Give specific examples of networks with high bandwidth-delay products. How is the bandwidth-delay product used in relation to throughput efficiency?
- 4.2 Collect **three** set of measurements (run this command 3 times) measuring the bandwidth of the public iperf hosts in Table 1 and find the mean bandwidth for each host.
- 4.3 Take the mean bandwidth and calculate the bandwidth-delay product in kilobits. You may use the mean round-trip delay time from your ping experiments to use as the delay time. Plot a bar chart for each host showing your results. You may wish to use a logarithmic scale, if appropriate.

Explain your results making reference to your networking environment which you performed your measurements. How do your results reflect upon your actual internet link speed? Are there outliers in your data?
- 4.4 Plot the bandwidth-delay product versus the hop count, this time omitting any outliers. Is there any correlation? If not, why not?
- 4.5 For each of the three sections in the project, give a brief description on your methodology, including all commands invoked, plotting software and operating system used, networking conditions and bandwidth of your internet link that will allow others to analyse your results.

Were there any variables which may affected the accuracy or reliability of your results? How might you improve upon these?

Reflect on the ease/difficulty on collecting accurate information in networks, in general. What are some aspects that hinder or assist this process?

5 Administration

This project can be performed individually or in pairs (2 people max) and is worth 10% of your overall mark in the subject. We will issue further information regarding groupwork in the weeks to come, including the submission of an individual reflection report.

5.1 Getting help

If you have any questions, the LMS discussion board will be a useful resource in resolving any issues. If your concern is a personal matter, then you should email the subject coordinator.

Any answers posted by the subject coordinator or the academic staff on the LMS discussion board will be considered as part of the project specification. Any announcements made about the assignment in the lectures will be considered part of the project specification. In addition, please keep an eye on any LMS announcements to any changes made to the project specification.

5.2 Academic honesty

Turnitin will be used as a deterrent against academic dishonesty. All gathered results, plots and discussion must be your own original work. All other information must be cited in an academically recognised style.

Plagiarism and collusion constitutes cheating. Disciplinary action will be taken against students who engage in plagiarism and collusion in accordance with university statute 13.1.18. For examples and more information on what constitutes plagiarism and collusion, please visit <http://academichonesty.unimelb.edu.au>.

5.3 Report submission

The deadline for the final report submission is as specified at the start of this project document.

The report will consist of all relevant discussion, graphs, data and answers from the experimentation conducted in this project. You must place the raw data as screens or copy paste them to the Appendix of the document, however, the diagrams like charts, flow diagrams (if any) and so on relevant to the discussion will be placed in the main document and not the appendix. Every diagram and/or raw measure used for a specific question must be referred to for us to verify the result. All plots and figures must be appropriately labelled. Any information obtained that is not of your own work must be cited.

The report must be submitted as a PDF file via Turnitin on the LMS. Please include as the first page, a cover sheet containing your name and login user name. The report is to be formatted on A4 sized paper in 10 pt text, 1.5 line spacing and a minimum 3.8 cm margins on the left and right sides.